


Agenda:

- Intro to sequence modeling
- Brief overview of RNNs
- Overview of Transformers and self Attention
- Multinead Attention
- Revisiting Tensor model parallelism & megatron-LM


*Intro to Sequence Modeling

- our new powerful LMs can:
 - 1) summarize large volumes of text
 - 2) Answer questions about world
 - 3) carry out long conversations
- in general we can think of these tasks as "sequence modeling" - mapping input sequences to output sequences:

- can also be useful for: audio, video, images, ekg/mri signals, data from IoT sensors; not just text!



* Today: sequence models for language



↳ model needs:

- world knowledge
- common sense
- (english)
- linguistic knowledge

* A LM is a compact probability distribution over sequences of tokens

→ model has access to vocab w/ tokens $x_1 \dots x_n$

→ model assigns sequence of tokens

some probability in 0..1

→ suppose vocab is: [lights, off, turn, the]:

$$P("turn off the lights") \Rightarrow .03$$

$$P("off turn the lights") \Rightarrow .002$$

... and so on

* we could model a joint distribution:

- choose some length L

- have LM learn joint distribution of tokens over variables x_1, x_2, \dots, x_n

$\rightarrow p(\text{"I like to eat cake"})$


$\rightarrow P(\text{"The movie was very funny"})$

\geq_{FRL}^S

* BUT this has 3 cons:

- 1) Ineffective - we typically want to sample sequences of different length
- 2) TOO many options - as length increases, # of possible sequences grow (hard to represent / space)
- 3) Imprecise: probabilities will represent sequences "on average" \rightarrow not best overall sequence


* Another Approach: RNNs - sota in 2016



- idea: capture "history" of all previously seen tokens to predict next token
- update HIDDEN STATE ($a_1, a_2 \dots a_{t+1}$):
 - use H.S. to generate output tokens

* challenge 1: modeling long-range dependencies

"The
counselor
helped
frome
the
situation"



"The" has
gone through
many iterations
by getting to
situation

* challenge 2: difficult to train!

- recall: backprop needs to go through many layers

- easy to get into a situation where model becomes **UNSTABLE** or **STOPS learning**

gradients
too large


gradients
too small

* challenge 3: FW / BW passes CAN'T be parallelized:

- each timestep needs to be processed before we move onto next step

* RNNs + Attention made some progress


- idea: all tokens should interact w/ representation of other tokens (helps w/ challenge 1)



* current mainstream approach:
TRANSFORMERS



- 1) uses attention idea
- 2) parallelizable
- 3) In last 7 years - we've gotten better at training




* from original google paper on transformers - google et al.

Figure 1: The Transformer - model architecture.

*Before we get into specifics, we can use transformers in a few diff. ways:

- Encoder only



example:

BERT

- purpose: use final representation to perform classification (e.g., sentiment analysis) or get good "representations" of sequence → create embeddings of a document
- processes text bidirectionally
- training objective: span prediction - mark some words so model can predict


- Decoder only

- purpose: use final representation to generate text

- training objective: next token prediction

- processes text in one direction (from left to

(right):



example:
GPT

• or both (encoder-decoder):



* Building Intuition for Attention:


- The meaning of "frame" depends on context:
 - "The counselor helped frame the situation"
 - "I hung up the photo frame"
- Attention op computes how much each token is influenced by other tokens:

The agreement on the European Economic Area was signed in August 1992.
L'accord sur la zone économique européenne a été signé en août 1992.
<end>

Il convient de noter que l'environnement marin est le moins connu de l'environnement. Il doit être noté que le milieu marin est le moins connu des milieux environnementaux.
It should be noted that the marine environment is the least known of environments.
<end>

Bandanau et al (Neural machine translation...)

Attention is All you need paper



Vaswani et al., 2017

*Let's try to break down self attention

→ given a sequence $(x_1, x_2 \dots x_n)$

"the counselor helped frame the

situation"

→ for each item x_i , compute how

much it should "pay attention" to

each value in the sequence


→ To do this, for each token x_i :

→ treat it as a "query token"

→ map entire sequence to "keys" w/
corresponding "values"

→ for query token, find all keys to
pay attention to → and take
those keys values

cs229s slides



$$x_i^{\text{NEW}} = 0.25 v_2 + 0.45 v_i + 0.3 v_N$$

J

create new version
of x_i by combining
values at some
probability

- probabilities
calculated w/ query
and all keys

* Full Attention Algorithm:

1) Transform each token x_i to set:
(multiply by
specific
pre-trained
matrix)

q_i k_i v_i
query vector key vector value vector

2) Compute dot product between q_i and
all k_i 's in sequence:

$$O = [q_i \cdot k_1, q_i \cdot k_2, q_i \cdot k_3, \dots q_i \cdot k_N]$$

3) Pass O through softmax to set probabilities

$$S = \text{Softmax}(O)$$

4) Take weighted sum of values given softmax output:

$$N$$


$$x_i^{\text{new}} = \sum_{l=1}^N S_l \cdot v_l$$

* This is great! We can compute in a batch:

N = sequence length / #of tokens in model input

d = model dimension

→ given sequence $(x_1 \dots x_n)$, model will get it as an $N \times d$ matrix of embeddings




* Attention step 1: compute all queries, keys, values

→ use matrix multiply to get

query (Q), key (K) and value (V)
matrices

learnable params



*2) Now how do we perform "lookup" and match queries to keys?

$$Q = Q K^T$$

Matrix multiply:

Note:

- Distance b/wn any

two tokens is now

$O(1)$ vs. $O(N)$ in iterating through KNN

- Parallelizable: can compute pairwise interactions in parallel

*3) How do we pick the "best" key



matches for our query?

→ Dot product

results in "score


b/wn each token pair

$\{i, j\}$




→ now for each query token i , we **NORMALIZE** dot product scores so they sum to 1

→ scores \sim weights that reflect how much query " i " matches w/ key " j "




→ now : take **WEIGHTED** sum of relevant tokens , where weights are these attention weights:



* What is MULTİHEAD ATTENTION?

- each head operates on smaller # of dimensions
- if input is d -dimensional, and we have h heads, each attention op. is over d/h inputs

Is 16 heads better than one
article



* where are we?






Figure 1: The Transformer - model architecture.

* ADD Feed-forward networks

→ so far: we haven't applied non-linearities;

just taken weighted averages

non-linearities improve expressivity



* use standard techniques to make architectures

stable:

1) Residual connections:

$$\text{Layer output} = \text{Attention output} + \text{Att. input}$$

2) Layer normalization: normalize output to

have zero mean / std 1 to keep scale

manageable

3) scale dot product: divide attention

weights by sqrt of model dimensions

are too big otherwise:

$$\text{Att. weights} = \frac{\text{softmax}(QK^T)}{\sqrt{d}}$$

Handling Position

→ problem: we don't know which token is in each position in our encoding!

The food was **good**. It had enough flavor and was not too expensive. However, the service was **poor**.

Good refers to food and **poor** refers to service.


However,
food had
was expensive.
It was good.
and flavor service
not the poor enough

* solution: Add Position Encodings:

(before passing input through).

$$x_i = t_i + p_i$$

↳ could also concatenate
in practice they are added



* output layers: (decoder)


- predict "next token" from full vocab size ✓

using output repr

1) output layer w learned map from
 $d \rightarrow V$ dimension.

↳ values associated w/ each vocab
↳ "logit"

- 2) Add softmax to get scores summing to 1
- 3) output token w/ HIGHEST softmax score



* revisiting Tensor model parallelism

in context of transformers (megatron!)

→ feed forward part might include:

- two-layer multi-layer perceptron (MLP)
- MLP = 2 matrix multiplies ($G \cdot F^{MN}$) +

GeLU non-linearity

$$Y = \text{GeLU}(X \cdot A)$$

$$Z = \text{Dropout}(Y)$$

- DISCUSS: should we split A
along columns or rows?

General lecture flow:

- Lecture 2 of CS229s 2023, given by Simran Arora: https://docs.google.com/presentation/d/1Pqu-TYLSnL9XNxF6BHIWahEtgNGNbmlwQzpgxduJJc/edit#slide=id.g2863f73f1c2_0_0
- Transformers and Attention Lecture of CMU 15-442

Images:

- Sequence Modeling: Simran's slides, # 7
- RNN breakdowns and challenges: Simran's slides, 34-37
- Attention Figure: Attention is All You need Paper (Vaswani et. Al): <https://arxiv.org/pdf/1706.03762.pdf>
- Encoder, Decoder, Encoder-Decoder diagrams: slides 27-29 of simran's lecture
- Heatmap of attention: Neural Machine Translation by Jointly Learning to Align and Translate: <https://arxiv.org/abs/1409.0473>
- Self attention matrix diagrams: 44-55 of Simran's lecture
- Multi head attention diagram: Are 16 heads better than 1: <https://blog.ml.cmu.edu/2020/03/20/are-sixteen-heads-really-better-than-one/>
- Transformers Architecture Build Up: Slide 66 - end of Simran's lecture

Other references:

Megatron Paper: <https://arxiv.org/pdf/2104.04473.pdf>